
ECE421
2021
Assignment
Five

Chase McDougall
Engineering Science
2023

Question 1

Fit_on_texts loops through the words in the given “messages” variable and updates the frequency tracking dictionary for each of the words it encounters. Either incrementing the count by 1 or initializing it to 1 if the word is contained within the dictionary keys or not respectively.

Question 2

Takes in a list of words and returns a list of the given words indices. These indices correspond to their placement within the frequency dictionary, where smaller indices correspond to higher frequency elements.

Question 3

Pad_sequences takes in a list of sequences and adds padding to the beginning of the sequence such that all sequences in the returned list are of the same length. This will be the length of the longest sequence unless a length limit is passed into the function. The default value for padding is 0 and the 0th index in the frequency dictionary is reserved for padding, though the padding value can be adjusted.

Question 4

We have a 2D array where each element of the array is a sequence of integers. These integers do not exceed 1999 which is the expected bound of sequence due to our set variables. We see that the shorter sequences have been padded with 0s at the beginning of the sequence. Therefore the acquired array for messages_train is of the form we expected.

Question 5

Depends on our definition of a “sentence”:

Since we cannot initiate or call using a string as input, we will assume that by “sentence” we mean a sequence of integers representing word indices.

However, even this is not an acceptable input for __init__ of the class.

Therefore the __init__ function will not work if a “sentence” is passed in, as it expects only initial integer inputs.

On initialization the class takes in two integer values which are used to define the dimensions of the initial weight matrix.

Calling the class function will use our input for indexing our initial weight matrix.

Our input sequence will be used to return a new matrix as follows:

Each element of the sequence will correspond to a row in the output matrix.

If the i^{th} element in this sequence is equal to k , then the i^{th} row in the output will be equal to the k^{th} row of the initial weight matrix. If k exceeds the maximum index of the initial weight matrix, then the corresponding output row will be equal to the last row of the initial weight matrix.

Question 6

```
update_gate =
objax.functional.sigmoid(jn.dot(update_w,x)+jn.dot(update_u,state)+update_
b)

    # fill this in r_t
    reset_gate =
objax.functional.sigmoid(jn.dot(reset_w,x)+jn.dot(reset_u,state)+reset_b)

    # fill this in h_hat_t
    output_gate = objax.functional.tanh(jn.dot(output_w, x) +
jn.dot(output_u, jn.multiply(reset_gate, state)) + output_b)
```

Question 7

```
def cumsum_norm(nums):
    # Calculate Cumulative Sum of nums in array
    return_array = []
    sum = 0
    for num in nums:
        sum += num
        return_array.append(sum)
    return jn.array(return_array)
x = jn.array([1,2,3,4,5])
print('Normal Cumulative Sum:')
print(cumsum_norm(x))

# Now Using lax.scan:
def cumsum_lax(sum, num):
    new_sum = sum + num
    return new_sum, new_sum

output, results = lax.scan(cumsum_lax, 0, x)
print('Lax Cumulative Sum:')
print(results)
```

Output:

```
Normal Cumulative Sum: [1 3 6 10 15]
Lax Cumulative Sum:    [1 3 6 10 15]
```

[lax.scan](#):

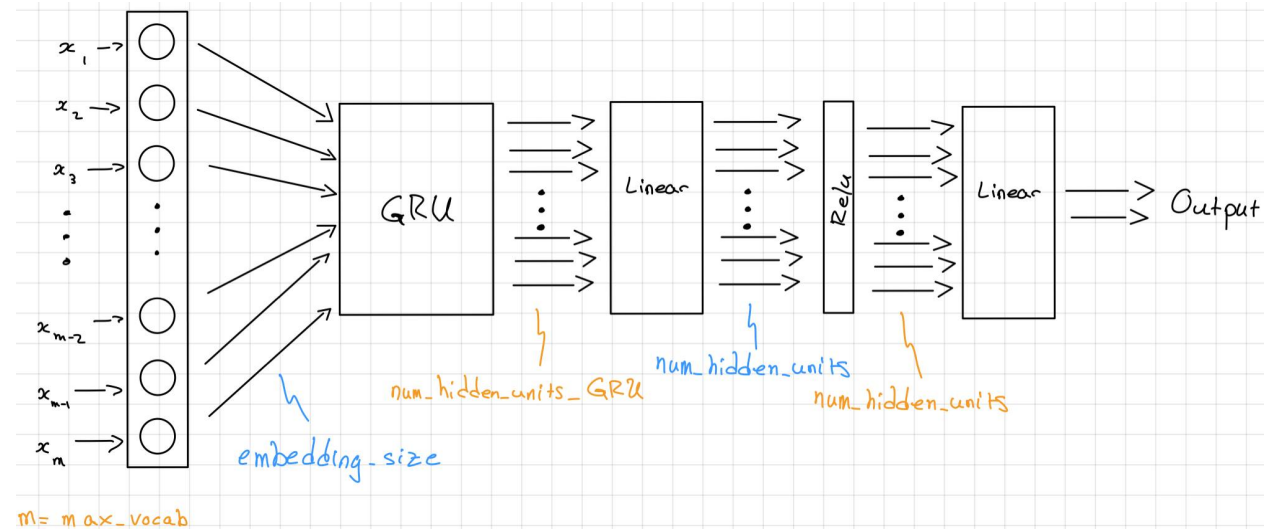
```
def scan(f, init, xs, length=None):  
    if xs is None:  
        xs = [None] * length  
    carry = init  
    ys = []  
    for x in xs:  
        carry, y = f(carry, x)  
        ys.append(y)  
    return carry, np.stack(ys)
```

The `lax.scan` function takes in a minimum of three inputs: a function, an initial value, and an array. The initial value will be used as our starting carry value and the preceding operations will manipulate it during the process. The array `x` will be iterated through in the passed in function augmenting the provided initial value.

In the given example of a cumulative sum, we see that the `lax.scan` function can be used to replace a traditional Python loop for improved performance and a reduction in written code. We will notice that the function we pass into `lax.scan` needs to take in two inputs as well as return two outputs. These inputs represent the current carry as well as the next value from the input array. In the case of the cumulative sum these two values are simply added together and their sum is returned as both the new carry and value at that point in time to be appended to our result array.

Question 8

```
gru_rnn = objax.nn.Sequential([
    Embed(max_vocab, embedding_size),
    GRU(embedding_size, num_hidden_units_GRU),
    objax.nn.Linear(num_hidden_units_GRU, num_hidden_units),
    objax.functional.relu,
    objax.nn.Linear(num_hidden_units, 2)
])
```

Question 9**Question 10**

```
opt = objax.optimizer.SGD(gru_rnn.vars())
```

Question 11

```
def train(EPOCHS = 20, BATCH = 20, LEARNING_RATE = 9e-4):
    avg_train_loss_epoch = []
    avg_val_loss_epoch = []
    train_acc_epoch = []
    val_acc_epoch = []

    for epoch in range(EPOCHS):
        avg_train_loss = 0 # (averaged) training loss per batch
        avg_val_loss = 0 # (averaged) validation loss per batch
        train_acc = 0 # training accuracy per batch
        val_acc = 0 # validation accuracy per batch

        # shuffle the examples prior to training to remove correlation
        train_indices = np.arange(len(messages_train))
        np.random.shuffle(train_indices)
        for it in range(0, messages_train.shape[0], BATCH):
            batch = train_indices[it:it+BATCH]
            avg_train_loss += float(train_op(messages_train[batch],
labels_train[batch], LEARNING_RATE)[0]) * len(batch)
            train_prediction = predict(messages_train[batch]).argmax(1)
            train_acc += (np.array(train_prediction).flatten() ==
labels_train[batch]).sum()
            train_acc_epoch.append(train_acc/messages_train.shape[0])
            avg_train_loss_epoch.append(avg_train_loss/messages_train.shape[0])

        # run validation
        val_indices = np.arange(len(messages_valid))
        np.random.shuffle(val_indices)
        for it in range(0, messages_valid.shape[0], BATCH):
            batch = val_indices[it:it+BATCH]
            avg_val_loss += float(loss_function(messages_valid[batch],
labels_valid[batch])) * len(batch)
            val_prediction = predict(messages_valid[batch]).argmax(1)
            val_acc += (np.array(val_prediction).flatten() ==
labels_valid[batch]).sum()
            val_acc_epoch.append(val_acc/messages_valid.shape[0])
            avg_val_loss_epoch.append(avg_val_loss/messages_valid.shape[0])
```

```
    print('Epoch %04d  Training Loss %.2f Validation Loss %.2f Training
Accuracy %.2f Validation Accuracy %.2f' % (epoch + 1,
avg_train_loss/messages_train.shape[0],
avg_val_loss/messages_valid.shape[0],
100*train_acc/messages_train.shape[0],
100*val_acc/messages_valid.shape[0]))

# Print Test Accuracy
print(f"Test Accuracy: {accuracy(test_data)}")

#Plot training loss
plt.title("Train vs Validation Loss")
plt.plot(avg_train_loss_epoch, label="Train")
plt.plot(avg_val_loss_epoch, label="Validation")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend(loc='best')
plt.show()

plt.title("Train vs Validation Accuracy")
plt.plot(train_acc_epoch, label="Train")
plt.plot(val_acc_epoch, label="Validation")
plt.xlabel("Epoch")
plt.ylabel("Accuracy (%)")
plt.legend(loc='best')
plt.show()
```


Question 12

Using SGD.

In the following table we have the accuracy after training for 20 epochs.

Validation Accuracy	Test Accuracy
81.40%	82.16%

The following Figure depicts the Training vs Validation accuracies during training over 20 epochs using SGD.



We have a relatively smaller generalization gap of $82.16\% - 81.40\% = 0.76\%$, which means that the model generalizes well to unseen data. Especially considering that the larger of the two accuracies is from the accuracy on the test set. I.e. unseen data.

Question 13

Reused code from previous parts with the exception of changing opt to:

```
opt2 = objax.optimizer.Adam(gru_rnn2.vars())
```



Question 14

Using Adam.

In the following table we have the accuracy after training for 20 epochs.

Validation Accuracy	Test Accuracy
88.30%	87.28%

The following Figure depicts the Training vs Validation accuracies during training over 20 epochs using Adam.



We now see a generalization gap of 1.02% with Validation accuracy now being greater than Test accuracy. While this is still a relatively small gap meaning our model generalizes well, it is now in the opposite direction of what we found with SGD. Furthermore, it is worth noting that while we see large fluctuations in accuracy with SGD, these fluctuations are not as present in our Adam training model. Specifically we see that the Validation remains relatively constant, while only the training accuracy climbs. Looking at a single epoch we obtain the following:

Validation Accuracy	Test Accuracy
85.74%	85.70%

It is interesting that when using Adam instead of SGD we are able to see such large accuracy values at such low numbers of epochs. It is also interesting to note that as we train the model over more epochs we see an increase in the generalization gap as it goes from roughly 0.04% to 1.02%.

Question 15

The following formulas illustrate how the Adam optimizer makes adjustments to our w parameters, with the last equation being the overall update step and other equations being used within it.

$$\begin{aligned}v_k &= \beta_1 v_{k-1} + (1 - \beta_1) \nabla f(.; w_{k-1}) \\s_k &= \beta_2 s_{k-1} + (1 - \beta_2) (\nabla f(.; w_{k-1}))^2 \\ \hat{v}_k &= \frac{v_k}{1 - \beta_1^k} \\ \hat{s}_k &= \frac{s_k}{1 - \beta_2^k} \\ w_k &= w_{k-1} - \eta \frac{\hat{v}_k}{\sqrt{\hat{s}_k} + \epsilon}\end{aligned}$$

Alternatively the following represents the update approach of SGD:

$$w_k = w_{k-1} - \eta \nabla f(.; w_{k-1})$$

We see that both options follow similar conventions with only the term following η differing between them.

Adam is said to be an extension of SGD which combines the advantages of two other advantages:

- Adaptive Gradient Algorithm (AdaGrad)
Maintains a per-parameter learning rate that improves performance on problems with sparse gradients (e.g. natural language and computer vision problems).
- Root Mean Square Propagation (RMSProp)
Maintains per-parameter learning rates that are adapted based on the average of recent magnitudes of the gradients for the weight (e.g. how quickly it is changing). This means the algorithm does well on online and non-stationary problems (e.g. noisy).

[1]

Adam maintains the benefits of both of these by not only adapting learning rates based on the average first moment, but also incorporating the average second moment of the gradient.

Question 16

Early Stopping Code:

```
if (avg_val_loss_epoch[len(avg_val_loss_epoch)-1] >
    avg_val_loss_epoch[len(avg_val_loss_epoch)-2]):
    if loss_increase:
        loss_seq += 1
    else:
        loss_increase = True
        loss_seq = 1

    if (loss_seq >= max_patience_window):
        print('Epoch %04d  Training Loss %.2f Validation Loss %.2f
Training Accuracy %.2f Validation Accuracy %.2f' % (epoch + 1,
avg_train_loss/messages_train.shape[0],
avg_val_loss/messages_valid.shape[0],
100*train_acc/messages_train.shape[0],
100*val_acc/messages_valid.shape[0]))
        print("Validation Error Increased. BREAK!")
        break
else:
    loss_increase = False
```

Validation Accuracy	Test Accuracy
87.96% at Epoch 4	87.76% at Epoch 10

Using the Adam Optimizer and a patience window of 5 epochs we get the accuracies shown in the above table. The early stopping algorithm stopped training at Epoch 10 after 5 epochs of increasing Validation Loss.



References

1. J. Brownlee, "Gentle introduction to the adam optimization algorithm for deep learning," *Machine Learning Mastery*, 12-Jan-2021. [Online]. Available: <https://machinelearningmastery.com/adam-optimization-algorithm-for-deep-learning/>. [Accessed: 30-Nov-2021].
-

Appendix

Question 11/12 Training Info and Plots:

Epoch 0001	Training Loss	138.15	Validation Loss	137.26	Training Accuracy	55.02	Validation Accuracy	56.98
Epoch 0002	Training Loss	134.92	Validation Loss	131.05	Training Accuracy	60.40	Validation Accuracy	59.58
Epoch 0003	Training Loss	129.35	Validation Loss	123.64	Training Accuracy	64.07	Validation Accuracy	66.42
Epoch 0004	Training Loss	122.83	Validation Loss	117.52	Training Accuracy	67.57	Validation Accuracy	68.12
Epoch 0005	Training Loss	118.88	Validation Loss	112.06	Training Accuracy	69.70	Validation Accuracy	70.72
Epoch 0006	Training Loss	114.04	Validation Loss	106.69	Training Accuracy	71.46	Validation Accuracy	73.84
Epoch 0007	Training Loss	110.64	Validation Loss	96.85	Training Accuracy	73.32	Validation Accuracy	77.06
Epoch 0008	Training Loss	107.38	Validation Loss	96.30	Training Accuracy	74.41	Validation Accuracy	77.82
Epoch 0009	Training Loss	106.11	Validation Loss	101.11	Training Accuracy	74.98	Validation Accuracy	74.88
Epoch 0010	Training Loss	100.92	Validation Loss	170.34	Training Accuracy	76.47	Validation Accuracy	61.32
Epoch 0011	Training Loss	99.80	Validation Loss	191.08	Training Accuracy	77.36	Validation Accuracy	55.92
Epoch 0012	Training Loss	109.32	Validation Loss	134.68	Training Accuracy	73.24	Validation Accuracy	56.04
Epoch 0013	Training Loss	132.79	Validation Loss	134.45	Training Accuracy	61.35	Validation Accuracy	56.66
Epoch 0014	Training Loss	126.89	Validation Loss	117.93	Training Accuracy	65.43	Validation Accuracy	69.90
Epoch 0015	Training Loss	122.21	Validation Loss	112.42	Training Accuracy	68.46	Validation Accuracy	72.50
Epoch 0016	Training Loss	114.38	Validation Loss	103.24	Training Accuracy	71.59	Validation Accuracy	75.20
Epoch 0017	Training Loss	103.94	Validation Loss	111.68	Training Accuracy	75.30	Validation Accuracy	72.14
Epoch 0018	Training Loss	95.55	Validation Loss	84.72	Training Accuracy	77.88	Validation Accuracy	80.50
Epoch 0019	Training Loss	91.60	Validation Loss	91.84	Training Accuracy	79.44	Validation Accuracy	77.64
Epoch 0020	Training Loss	87.34	Validation Loss	85.57	Training Accuracy	80.40	Validation Accuracy	81.40





Question 14 Training Info and Plots:

```

Epoch 0001 Training Loss 90.10 Validation Loss 62.42 Training Accuracy 78.83 Validation Accuracy 86.56
Epoch 0002 Training Loss 57.72 Validation Loss 57.53 Training Accuracy 88.71 Validation Accuracy 87.30
Epoch 0003 Training Loss 53.50 Validation Loss 56.69 Training Accuracy 89.52 Validation Accuracy 87.98
Epoch 0004 Training Loss 50.88 Validation Loss 55.56 Training Accuracy 90.16 Validation Accuracy 87.98
Epoch 0005 Training Loss 48.69 Validation Loss 54.65 Training Accuracy 90.66 Validation Accuracy 88.50
Epoch 0006 Training Loss 46.78 Validation Loss 57.71 Training Accuracy 91.05 Validation Accuracy 88.08
Epoch 0007 Training Loss 44.27 Validation Loss 56.88 Training Accuracy 91.60 Validation Accuracy 87.84
Epoch 0008 Training Loss 42.96 Validation Loss 57.95 Training Accuracy 91.89 Validation Accuracy 88.12
Epoch 0009 Training Loss 40.98 Validation Loss 58.45 Training Accuracy 92.34 Validation Accuracy 87.94
Epoch 0010 Training Loss 39.28 Validation Loss 58.33 Training Accuracy 92.71 Validation Accuracy 88.30
Epoch 0011 Training Loss 36.97 Validation Loss 58.21 Training Accuracy 93.18 Validation Accuracy 88.20
Epoch 0012 Training Loss 34.20 Validation Loss 60.24 Training Accuracy 93.83 Validation Accuracy 88.60
Epoch 0013 Training Loss 32.12 Validation Loss 62.67 Training Accuracy 94.25 Validation Accuracy 88.50
Epoch 0014 Training Loss 30.77 Validation Loss 63.10 Training Accuracy 94.58 Validation Accuracy 88.74
Epoch 0015 Training Loss 28.02 Validation Loss 70.31 Training Accuracy 95.17 Validation Accuracy 88.56
Epoch 0016 Training Loss 25.94 Validation Loss 68.84 Training Accuracy 95.54 Validation Accuracy 88.28
Epoch 0017 Training Loss 24.85 Validation Loss 71.90 Training Accuracy 95.75 Validation Accuracy 88.16
Epoch 0018 Training Loss 22.72 Validation Loss 76.40 Training Accuracy 96.22 Validation Accuracy 88.44
Epoch 0019 Training Loss 21.27 Validation Loss 85.13 Training Accuracy 96.58 Validation Accuracy 87.80
Epoch 0020 Training Loss 20.35 Validation Loss 84.01 Training Accuracy 96.76 Validation Accuracy 88.30
Test Accuracy: 87.28%

```



Early Stopping:

```
Epoch 0001 Training Loss 92.26 Validation Loss 60.74 Training Accuracy 77.94 Validation Accuracy 87.06
Epoch 0002 Training Loss 59.01 Validation Loss 58.38 Training Accuracy 88.47 Validation Accuracy 87.26
Epoch 0003 Training Loss 53.33 Validation Loss 58.06 Training Accuracy 89.46 Validation Accuracy 87.50
Epoch 0004 Training Loss 50.57 Validation Loss 56.18 Training Accuracy 90.25 Validation Accuracy 87.96
Epoch 0005 Training Loss 47.90 Validation Loss 55.65 Training Accuracy 90.68 Validation Accuracy 88.02
Epoch 0006 Training Loss 45.91 Validation Loss 56.63 Training Accuracy 91.16 Validation Accuracy 87.68
Epoch 0007 Training Loss 43.90 Validation Loss 57.14 Training Accuracy 91.72 Validation Accuracy 87.62
Epoch 0008 Training Loss 42.54 Validation Loss 59.16 Training Accuracy 92.02 Validation Accuracy 87.64
Epoch 0009 Training Loss 40.61 Validation Loss 59.40 Training Accuracy 92.48 Validation Accuracy 87.92
Epoch 0010 Training Loss 38.47 Validation Loss 60.62 Training Accuracy 92.85 Validation Accuracy 87.32
Validation Error Increased. BREAK!
Test Accuracy: 87.76%
```

